

ORACLE®

AST Specialisation and Partial Evaluation for Easy High-Performance Metaprogramming

1st Workshop on Meta-Programming Techniques and Reflection (META)

Chris Seaton
Research Manager
Oracle Labs
November 2016

Safe Harbor Statement

The following is intended to provide some insight into a line of research in Oracle Labs. It is intended for information purposes only, and may not be incorporated into any contract. It is not a commitment to deliver any material, code, or functionality, and should not be relied upon in making purchasing decisions. Oracle reserves the right to alter its development plans and practices at any time, and the development, release, and timing of any features or functionality described in connection with any Oracle product or service remains at the sole discretion of Oracle. Any views expressed in this presentation are my own and do not necessarily reflect the views of Oracle.

Outline

- We are using a novel combination of techniques to create high performance implementations of existing languages
 - Truffle: framework for writing AST interpreters in Java
 - Graal: new dynamic (JIT) compiler for the JVM that knows about Truffle
- We've found that this combination of tools is particularly useful for easy, pervasive, consistent, high-performance metaprogramming implementations
- We'll show why this is and what it looks like
- We'll suggest what properties from Truffle and Graal could be useful to make sure future language implementation systems have

Truffle and Graal



HotSpot





HotSpot



HotSpot

JIT






Truffle



Graal

HotSpot

JIT




Truffle

Truffle

Truffle

Graal



Truffle


Truffle

Truffle

Graal

Truffle for AST interpreters


$x + y * z$



load_local x
load_local y
load_local z
call *
call +

pushq %rbp
movq %rsp, %rbp
movq %rdi, -8(%rbp)
movq %rsi, -16(%rbp)
movq %rdx, -24(%rbp)
movq -16(%rbp), %rax
movl %eax, %edx
movq -24(%rbp), %rax
imull %edx, %eax
movq -8(%rbp), %rdx
addl %edx, %eax
popq %rbp
ret


$x + y * z$



load_local x
load_local y
load_local z
call *
call +


pushq %rbp
movq %rsp, %rbp
movq %rdi, -8(%rbp)
movq %rsi, -16(%rbp)
movq %rdx, -24(%rbp)
movq -16(%rbp), %rax
movl %eax, %edx
movq -24(%rbp), %rax
imull %edx, %eax
movq -8(%rbp), %rdx
addl %edx, %eax
popq %rbp
ret

$x + y * z$




load_local x
load_local y
load_local z
call *
call +

pushq %rbp
movq %rsp, %rbp
movq %rdi, -8(%rbp)
movq %rsi, -16(%rbp)
movq %rdx, -24(%rbp)
movq -16(%rbp), %rax
movl %eax, %edx
movq -24(%rbp), %rax
imull %edx, %eax
movq -8(%rbp), %rdx
addl %edx, %eax
popq %rbp
ret




AST Interpreter
Uninitialized Nodes

T. Würthinger, C. Wimmer, A. Wöß, L. Stadler, G. Duboscq, C. Humer, G. Richards, D. Simon, and M. Wolczko. One VM to rule them all. In Proceedings of Onward!, 2013.




AST Interpreter
Uninitialized Nodes

Node Rewriting for Profiling Feedback



T. Würthinger, C. Wimmer, A. Wöß, L. Stadler, G. Duboscq, C. Humer, G. Richards, D. Simon, and M. Wolczko. One VM to rule them all. In Proceedings of Onward!, 2013.


Graal for partial evaluation



AST Interpreter
Rewritten Nodes


Compiled Code

T. Würthinger, C. Wimmer, A. Wöß, L. Stadler, G. Duboscq, C. Humer, G. Richards, D. Simon, and M. Wolczko. One VM to rule them all. In Proceedings of Onward!, 2013.





codon.com/compilers-for-free

Presentation, by Tom Stuart, licensed under a Creative Commons Attribution ShareAlike 3.0





T. Würthinger, C. Wimmer, A. Wöß, L. Stadler, G. Duboscq, C. Humer, G. Richards, D. Simon, and M. Wolczko. One VM to rule them all. In Proceedings of Onward!, 2013.

Node Rewriting to Update Profiling Feedback




Recompilation using Partial Evaluation



T. Würthinger, C. Wimmer, A. Wöß, L. Stadler, G. Duboscq, C. Humer, G. Richards, D. Simon, and M. Wolczko. One VM to rule them all. In Proceedings of Onward!, 2013.

Metaprogramming in Ruby



```
# Conventional send  
object.method_name(arg1, arg2, ...)  
# Metaprogramming send  
object.send('method_name', arg1, arg2, ...)
```

```
operator = exclude_end? ? :< : :<=
value.send(operator, last)
```

```
send("decode_png_resample_#{bit_depth}bit_value")
```

```
def method_missing(method, *args)
  @encapsulated_value.send(method, *args)
end
```

```
def method_missing(name, *args)
  if Color.respond_to?(name)
    return Color.send(name, *args)
  end
end
```

```
eval(generated_template, variables)
```

```
object.instance_variable_get('@variable_name')
object.instance_variable_set('@variable_name', value)
```

```
def eql?(other)
  @hash.eql?(other.instance_variable_get(:@hash))
end
```

Foundational techniques

Caching

```
a = [1, 2, 3]
puts a[2]
```

```
h = {1=>a, 2=>b, 3=>c}
puts h[2]
```

Class	Method name	Method
Array	[]	Array#[]
Hash	[]	Hash#[]
.... more entries ...		

one table per virtual machine, lots of entries

L. P. Deutsch and A. M. Schiffman. Efficient Implementation of the Smalltalk-80 System, 1984.

Inline caching

```
a = [1, 2, 3]
puts a[2]
```

Class	Method
Array	Array#[]

one table per call site, one entry

```
h = {1=>a, 2=>b, 3=>c}
puts h[2]
```

Class	Method
Hash	Hash#[]

one table per call site , one entry


Polymorphic inline caching

```
x = random(a, h)  
x[2]
```



one table per call site, one entry

```
x = random(a, h)  
x[2]
```



one table per call site , one entry

Polymorphic inline caching

U. Hözle, C. Chambers, and D. Ungar. Optimizing dynamically-typed object-oriented languages with polymorphic inline caches. In *ECOOP'91 European Conference on Object-Oriented Programming*, volume 512 of *Lecture Notes in Computer Science*. 1991.

```
x = random(a, h)  
x[2]
```

Class	Method
Array	Array#[]
Hash	Hash#[]
.... more entries ...	

one table per call site, multiple entries

```
x = random(a, h)  
x[2]
```

Class	Method
Array	Array#[]
Hash	Hash#[]
.... more entries ...	

one table per call site, multiple entries

Dispatch chains

S. Marr, C. Seaton, and S. Ducasse. Zero-overhead metaprogramming: Reflection and metaobject protocols fast and without compromises. In *Proceedings of the 36th ACM SIGPLAN Conference on Programming Language Design and Implementation*, 2015.

```
bit_depth = random(8, 16, 32)
send(image, "resample_#{bit_depth}bit")
```

Class	Method name	Method
Image	resample_8bit	Image#resample_8bit
Image	resample_16bit	Image#resample_16bit
Image	resample_32bit	Image#resample_32bit
.... more entries ...		

one table per call site, multiple entries

Why aren't these a solution on their own?

Caches are currently implemented manually

```
struct rb_call_cache {
    /* inline cache: keys */
    rb_serial_t method_state;
    rb_serial_t class_serial;

    /* inline cache: values */
    const rb_callable_method_entry_t *me;

    vm_call_handler call;

    union {
        unsigned int index; /* used by ivar */
        enum method_missing_reason method_missing_reason; /* used by method_missing */
        int inc_sp; /* used by cfunc */
    } aux;
};
```

You need somewhere to store an inline cache

```
a.foo(b)
```

```
a = [1, 2, 3]
```

```
a.sort
```


```
a.send(:foo, b)
```

You need somewhere to store an inline cache

a.foo(b)

a = [1, 2, 3]
a.sort

a.send(:foo, b)




a <= b

You need somewhere to store an inline cache

```
a.foo(b)
```

```
a = [1, 2, 3]  
a.sort
```

```
a.send(:foo, b)
```




You need somewhere to store an inline cache

```
@JRubyMethod(name = "send")
public static IRubyObject send(ThreadContext context, IRubyObject self, String name, IRubyObject[] args) {
    DynamicMethod method = searchMethod(name);
    return method.call(context, self, this, name, args);
}
```


You need somewhere to store an inline cache

```
@JRubyMethod(name = "sort")
public static IRubyObject sort(ThreadContext context, IRubyObject array, String name) {
    ...
    Arrays.sort(newValues, 0, length, new Comparator() {
        public int compare(Object o1, Object o2) {
            DynamicMethod method = searchMethod("<=>");
            return method.call(context, self, this, name, o1, o2);
        }
    });
    ...
}
```

Caches quickly become megamorphic



Caches quickly become megamorphic



How Truffle and Graal make a difference

An easy place to store state

```
class SendNode extends Node {  
    String methodName;  
    Node receiverNode;  
  
    public Object execute() {  
        Object receiver = receiverNode.execute();  
        Method method = receiver.lookup(methodName);  
        return method.call();  
    }  
}
```

An easy place to store state

```
class SendNode extends Node {  
    String methodName;  
    Node receiverNode;  
    Class cachedClass;  
    Method cachedMethod;  
  
    public Object execute() {  
        Object receiver = receiverNode.execute();  
        if (receiver.getClass() != cachedClass) {  
            cachedClass = receiver.getClass();  
            cachedMethod = receiver.lookup(methodName);  
        }  
        return cachedMethod.call();  
    }  
}
```

A DSL to write caches in just a couple of lines


```
@NodeChild("receiver")
class SendNode extends Node {
    String methodName;

    @Specialisation(guards = "receiver.getClass() == cachedClass")
    public Object execute(Object receiver,
                          @Cached("receiver.getClass()") Class cachedClass,
                          @Cached("receiver.lookup(methodName)") Method cachedMethod) {
        return method.call();
    }
}
```


A DSL to write caches in just a couple of lines

```
@NodeChildren({"receiver", "name"})
class SendNode extends Node {
    @Specialisation(guards = {"receiver.getClass() == cachedClass", "name.equals(cachedName)"})
    public Object execute(Object receiver,
                          String name,
                          @Cached("receiver.getClass()") Class cachedClass,
                          @Cached("name") String cachedName,
                          @Cached("receiver.lookup(name)") Method cachedMethod) {
        return method.call();
    }
}
```


Automatic splitting to push caches down the call stack



Automatic splitting to push caches down the call stack




Automatic splitting to push caches down the call stack




Results

```
def eql?(other)
  @hash.eql?(other.instance_variable_get(:@hash))
end
```


Relative performance of metaprogramming access to instance variables relative to conventional access



Slowdown of metaprogramming access to instance variables relative to JRuby+Truffle



Slowdown of Set#eq? relative to JRuby+Truffle



The important properties

Somewhere to store state

- Caching and profiling requires somewhere to store state
- Truffle's nodes are just Java objects, so you can store whatever you want in normal Java fields
- In Truffle you are almost always in a node, so you almost always have access to your state
 - Doesn't become inaccessible in compiled code

Low-effort caching

- Truffle's DSL makes it easy to add sophisticated polymorphic inline caches anywhere
- This is implemented using the state that we just mentioned
- Guards can be arbitrary Java expressions, or zero-overhead mutable flags using deoptimisation
- Supports an arbitrary number of guards

Dynamic optimisation

- Dynamic optimisation (JIT compilation) comes for free from Graal
- Partial evaluation removes degrees of freedom that aren't used
 - Allows us to add degrees of freedom to handle metaprogramming without worrying

Dynamic *deoptimisation*

- Allows us to make speculative optimisations and reverse them if they were wrong
- Allows functionality not used to be ‘turned off’ until it is needed
- Allows local variables to be lowered all the way to registers while still letting frames be accessed as if they were objects

Automatic inlining and splitting

- Removes the overhead of intermediate methods calls and indirection used in metaprogramming
- Allows state to be ‘pushed down’ the call stack to reduce polymorphism

Programmatic access to frames

- Allows local variables to be read and written from outside method activations
- Whole frames represented as objects
- Access to the list of frames currently on the stack

Conclusions

- We already knew how to make most (not all) of Ruby's metaprogramming functionality fast
- Existing mature Ruby implementations don't apply this knowledge
- Why? Because it was hard in practice to do it consistently and pervasively that they never got around to it

Conclusions

- Truffle and Graal make it so much easier
- We've identified what we think are the key properties that enable this
- I think Truffle and Graal are the only systems to provide effective implementations of these
- If you are implementing a metaprogramming language, use Truffle and Graal
- If you're making a new language implementation system, perhaps incorporate these same properties

Where to find more information

GitHub, Inc.

This organization Search

Pull requests Issues Gist

Search for 'github graalvm'

Graal Multi-Language VM
Next generation compilation technology supporting Java, Ruby, R, JavaScript, LLVM, and more.
<https://graalvm.github.io>

Repositories People 38 Teams 2

Filters Find a repository...

sulong Java ★ 211 ⚡ 19
Sulong, a dynamic runtime for LLVM-based languages.
Updated 6 minutes ago

graal-core Java ★ 122 ⚡ 33
Graal Compiler & Truffle Partial evaluator.
Updated 31 minutes ago

mx Python ★ 13 ⚡ 26

People 38 >


github.com/graalvm

The screenshot shows the GitHub organization page for 'graalvm'. The header includes the GitHub logo, navigation links for 'Pull requests', 'Issues', and 'Gist', and a search bar. A prominent orange search bar on the right contains the text 'Search for 'github graalvm''. The main content area displays three repositories: 'sulong' (Java, 211 stars), 'graal-core' (Java, 122 stars), and 'mx' (Python, 13 stars). Below the repositories is a 'People' section showing 38 members, many of whom have their standard GitHub profile pictures replaced by green and yellow pixelated versions. At the bottom right is a large orange button with the URL 'github.com/graalvm'.



Truffle and Graal: Fast Programming Languages With Modest Effort

Thursday, 14:20, Matterhorn 3 (this room)
SPLASH-I
Adam Welc



Acknowledgements

Oracle

Danilo Ansaloni
Stefan Anzinger
Cosmin Basca
Daniele Bonetta
Matthias Brantner
Petr Chalupa
Jürgen Christ
Laurent Daynès
Gilles Duboscq
Martin Entlicher
Brandon Fish
Bastian Hossbach
Christian Hummer
Mick Jordan
Vojin Jovanovic
Peter Kessler
David Leopoldseder
Kevin Menard
Jakub Podlešák
Aleksandar Prokopec
Tom Rodriguez

Oracle (continued)

Roland Schatz
Chris Seaton
Doug Simon
Štěpán Šindelář
Zbyněk Šlajchrt
Lukas Stadler
Codrut Stancu
Jan Štola
Jaroslav Tulach
Michael Van De Vanter
Adam Welc
Christian Wimmer
Christian Wirth
Paul Wögerer
Mario Wolczko
Andreas Wöß
Thomas Würthinger

Oracle Interns

Brian Belleville
Miguel Garcia
Shams Imam
Alexey Karyakin
Stephen Kell
Andreas Kunft
Volker Lanting
Gero Leinemann
Julian Lettner
Joe Nash
David Piorkowski
Gregor Richards
Robert Seilbeck
Rifat Shariyar

Alumni

Erik Eckstein
Michael Haupt
Christos Kotselidis
Hyunjin Lee
David Leibs
Chris Thalinger
Till Westmann

JKU Linz

Prof. Hanspeter Mössenböck
Benoit Daloze
Josef Eisl
Thomas Feichtinger
Matthias Grimmer
Christian Häubl
Josef Haider
Christian Huber
Stefan Marr
Manuel Rigger
Stefan Rumzucker
Bernhard Urban

University of Edinburgh

Christophe Dubach
Juan José Fumero Alfonso
Ranjeet Singh
Toomas Remmelg

LaBRI

Floréal Morandat

University of California, Irvine

Prof. Michael Franz
Gulfem Savrun Yeniceri
Wei Zhang

Purdue University

Prof. Jan Vitek
Tomas Kalibera
Petr Maj
Lei Zhao

T. U. Dortmund

Prof. Peter Marwedel
Helena Kotthaus
Ingo Korb

University of California, Davis

Prof. Duncan Temple Lang
Nicholas Ulle

University of Lugano, Switzerland

Prof. Walter Binder
Sun Haiyang
Yudi Zheng

Safe Harbor Statement

The preceding is intended to provide some insight into a line of research in Oracle Labs. It is intended for information purposes only, and may not be incorporated into any contract. It is not a commitment to deliver any material, code, or functionality, and should not be relied upon in making purchasing decisions. Oracle reserves the right to alter its development plans and practices at any time, and the development, release, and timing of any features or functionality described in connection with any Oracle product or service remains at the sole discretion of Oracle. Any views expressed in this presentation are my own and do not necessarily reflect the views of Oracle.

Integrated Cloud Applications & Platform Services

ORACLE®