

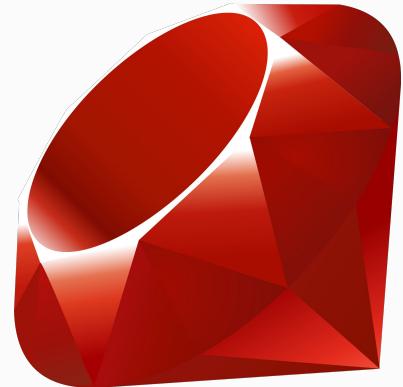
A History of Compiling Ruby

Chris Seaton

Senior Staff Engineer

Shopify

RubyConf 2021



What is a Ruby compiler?

- From Ruby code, to native machine code
- Can be via some intermediate system
- Can be ahead-of-time or just-in-time
- Requires some judgement and opinions to interpret, but I'm not trying to gate-keep



Did you know there have been at least...

25

...attempts to build a Ruby compiler?



Buy why?

- You can talk to anyone about compilers
- Compilers are really fun
- It's surprising how many there are... why has this happened?
- There's new interest in compilers for Ruby
- Lots of existing work to learn from - what's been tried before?
- We can trace compiler research through Ruby compilers
- Some of these projects were at risk of being lost to mists of time



What we're definitely not doing here...

- Talking about which is better or why
- Benchmarking or a competition
- Inevitably going to contain a few opinions



Compiler	Years active	Base VM	Stage	General approach	Frontend	Interpreter	Intermediate representations	Key authors
Hokstad	2008-present	Custom Ruby	AOT	Template compilation of an AST	Custom recursive descent and operator precedence parser	None	Enhanced AST	Hokstad
Hyperdrive	2019-2020	MRI	JIT	Tracing of YARV instructions then template compilation to Cranelift	Tracing YARV interpreter	Instrumented base interpreter	None	Matthews
IronRuby	2007-2011	Custom C#	JIT	Generation of CIL				Lam
JRuby	2006-present	Custom Java	JIT*	Generation of JVM bytecode	Parser to AST, to internal IR	Internal IR interpreter	CFG of linear RTL instructions	Nutter, Enebo, Sastry
LLRB	2017	MRI	JIT	Generation of LLVM IR				Kokubun
Ludicrous	2008-2009	MRI	JIT	Template compilation through DotGNU LibJIT				Brannan
MacRuby	2008-2013	MRI	AOT/JIT	Generation of LLVM IR				Sansonetti
MagLev	2008-2016	Custom Gemstone Smalltalk	JIT					McLain, Felgentreff
Natalie	2019-present	Custom C++	AOT	AST incrementally lowered to C++			Enhanced AST	Morgan
Ruby+OMR	2016-2017	MRI	JIT	Generation of J9 IR				Gaudet, Stodely
RTL MJIT	2017	MRI	JIT	Generation of C				Makarov
Rubinius	2008-2016	Custom C++ and Ruby	JIT	Generation of LLVM IR	Parser to AST, to custom stack bytecode	Stack bytecode	None	Phoenix, Bussink, Shirai
Rhizome	2017	MRI, JRuby, Rubinius	JIT	Conventional speculative compiler with in-process assembler	Base bytecode or IR to custom bytecode	Stack bytecode	Graphical sea-of-nodes	Seaton
RubyComp	2004		AOT					Alexandersson
RubyX	2014-2020		AOT	Conventional compiler with in-process assembler	Parser to AST	None	Multiple IRs gradually removing abstraction and lowering from AST to linear	Rüger
Ruby.NET	2008	Custom C#		Generation of CIL				Kelly
Rucy	2021		AOT	Template compilation to eBPF				Uchio
Sorbet Compiler	2019-present	MRI	AOT	Generation of MRI LLVM IR 'C' extension	Parser to AST	None	Sorbet's typechecking IR	Tarjan, Petrushko, Froyd
TenderJIT	2021	MRI	JIT	Lazy Basic Block Versioning with in-process assembler	Template compiler of YARV bytecode	Base interpreter	None	Patterson
Topaz	2012-2014	Custom RPython and Ruby	JIT	Metatracing of a stack bytecode interpreter	Parser to AST	Stack bytecode interpreter		Gaynor, Felgentreff
TruffleRuby	2013-present	Custom Java and Ruby	JIT	Partial evaluation of self-specialising AST	Parser to AST	Self-specialising AST interpreter	Graphical sea-of-nodes	Seaton, Daloze, Menard, Chalupa, MacGregor
XRuby	2006-2008	Custom Java	AOT	Template compilation to Java bytecode	Parser to AST	None	None	Zhi
yarv2llvm	2008-2010	MRI	JIT	Generation of LLVM IR				Hideki
YARV MJIT	2018-present	MRI	JIT	Generation of C		Base interpreter		Kokubun
YJIT	2020-present	MRI	JIT	Lazy Basic-Block Versioning with in-process assembler	Template compiler of YARV bytecode	Base interpreter	None	Chevalier-Boisvert

Key Ruby compilers today

JRuby	Ruby on the JVM
TruffleRuby	Ruby on the GraalVM with a core in Ruby
YARV MJIT	MRI 3x3's JIT
YJIT	Shopify's new MRI JIT
Sorbet Compiler	Stripe's Ruby to C extension compiler



Key Ruby compilers that passed us by

Rubinius Ruby “in Ruby” with LLVM

Topaz Ruby in RPython, like PyPy

IronRuby Ruby in C#

Ruby+OMR Ruby on J9



Fun niche compilers I'd recommend

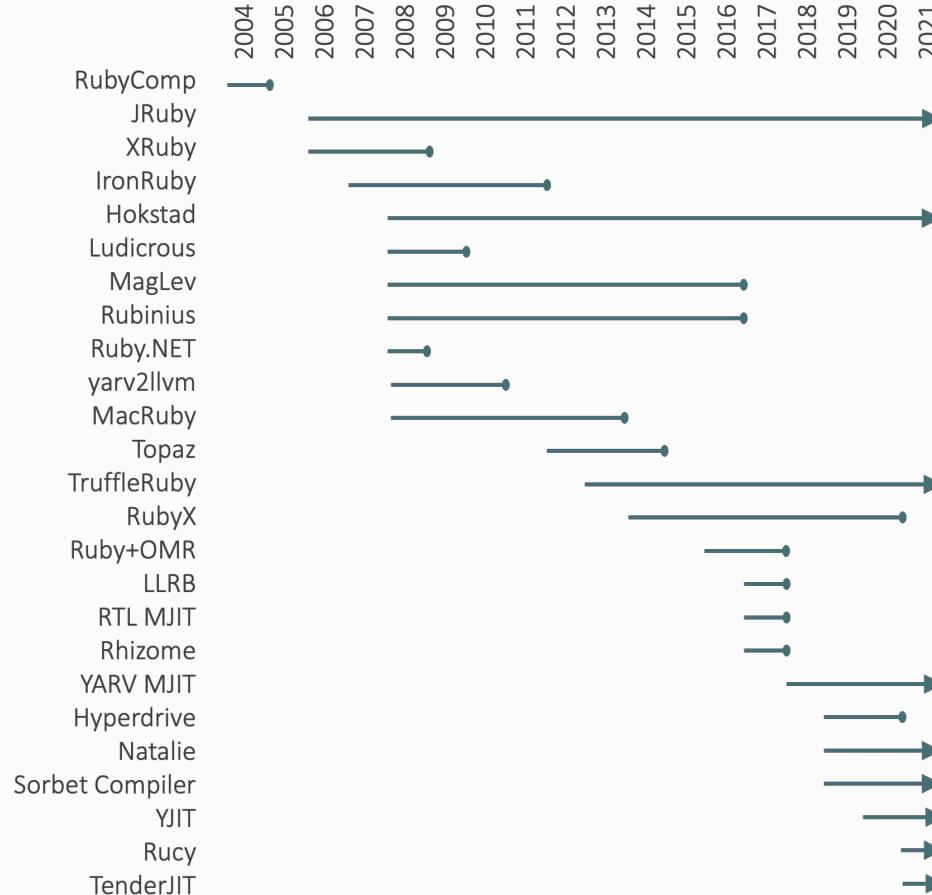
Hokstad Ruby directly to ARM that's easy to understand

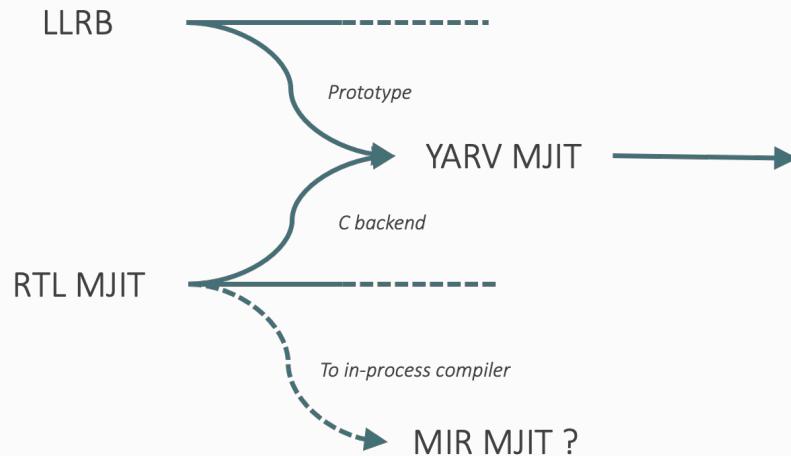
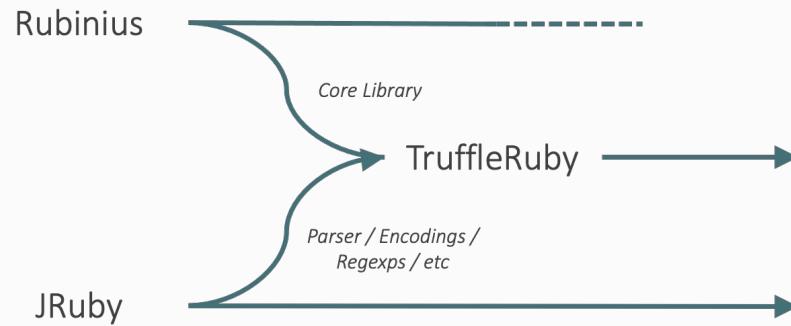
Hyperdrive Ruby with tracing

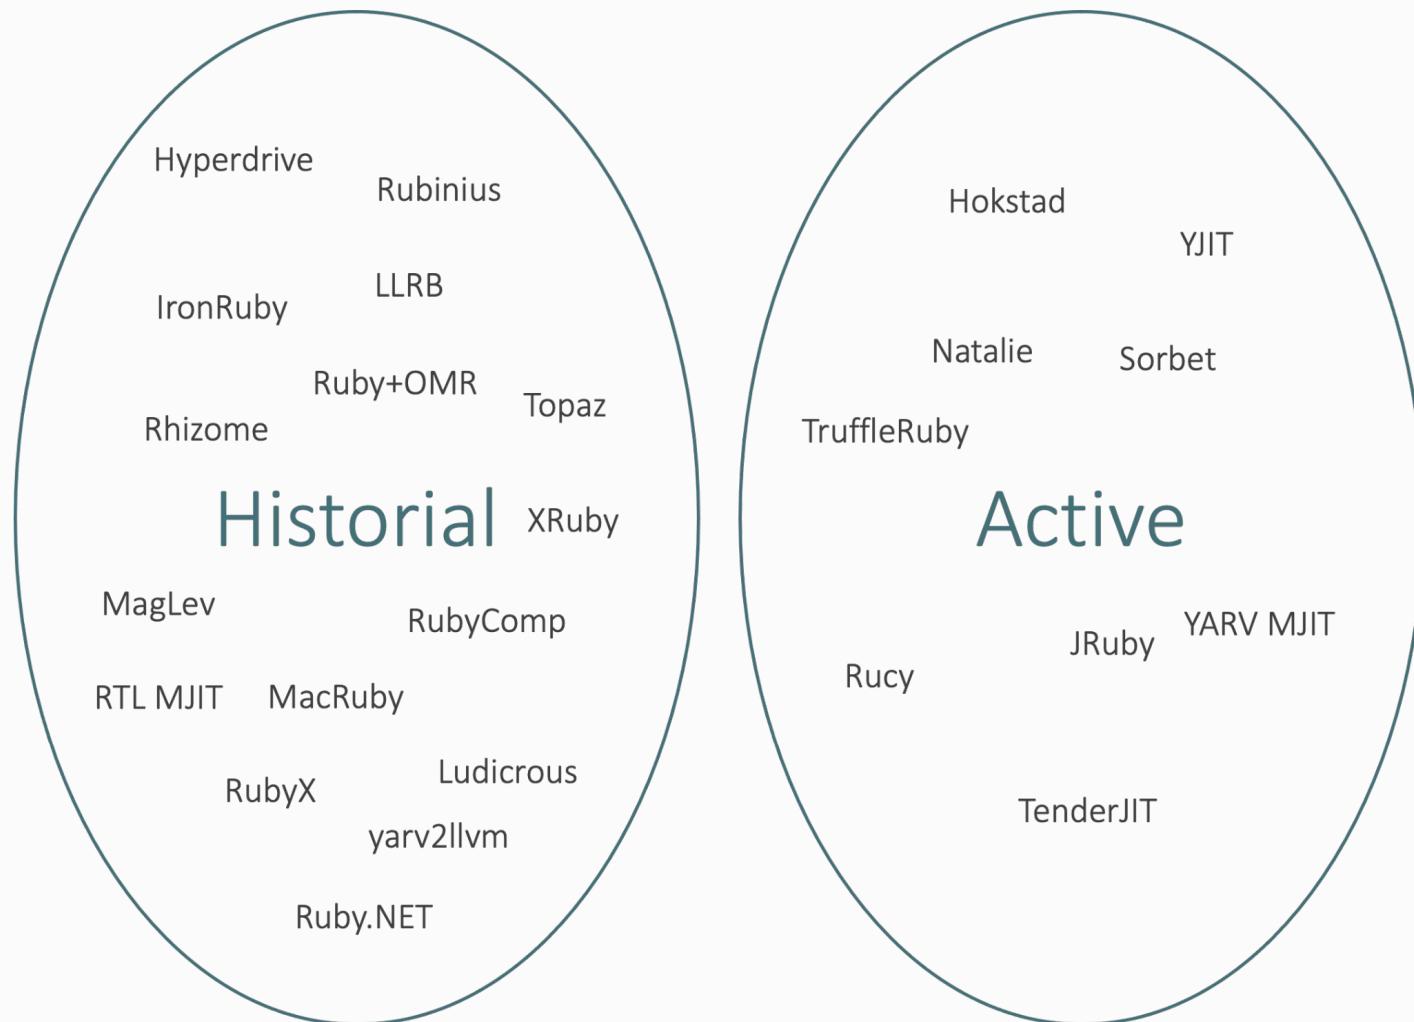
Natalie Ruby to C++ with great videos

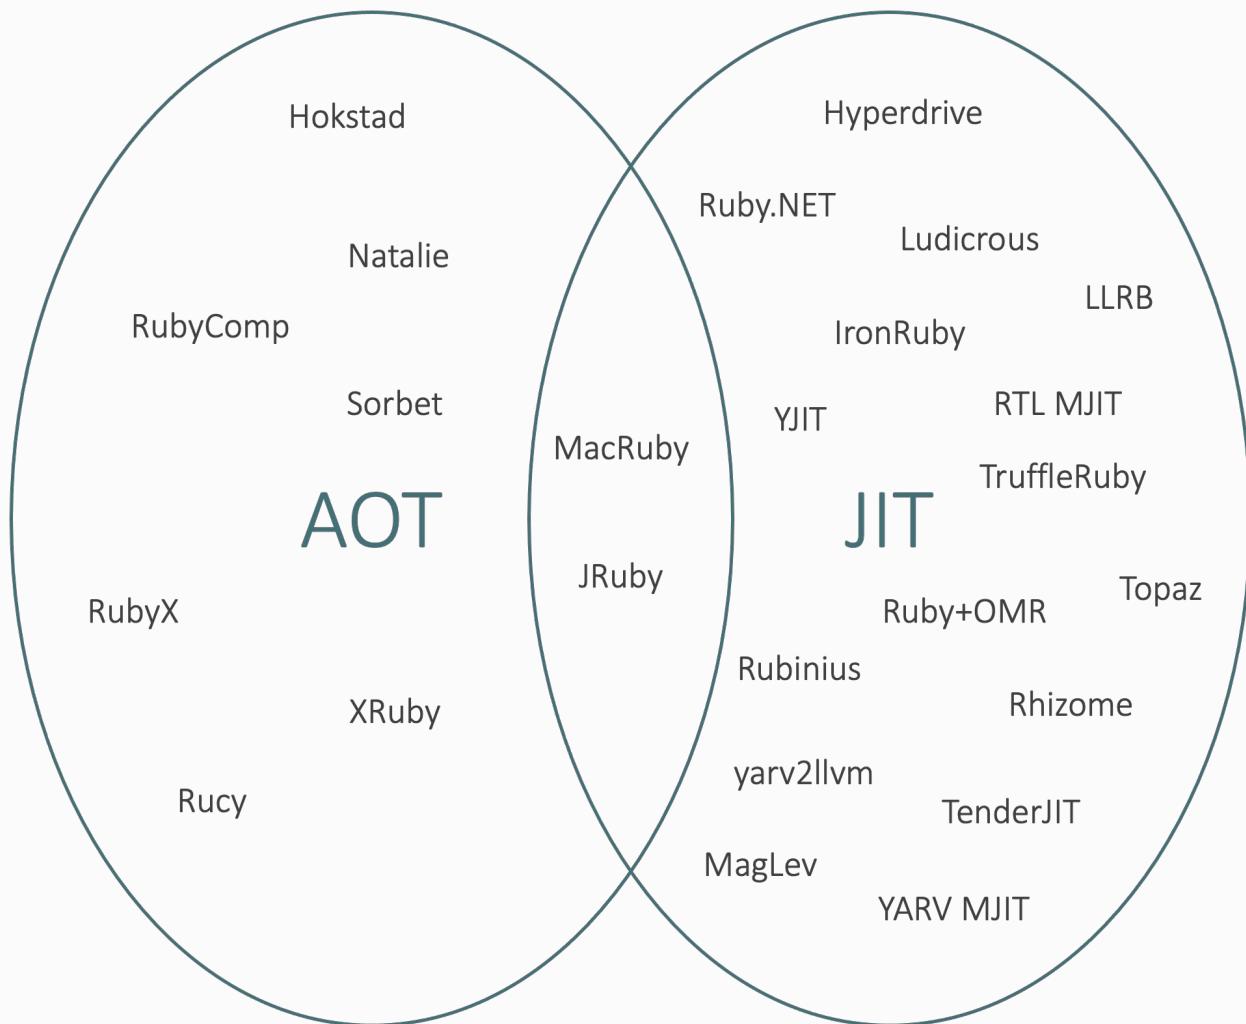
RubyX Ruby to ARM with a more advanced architecture

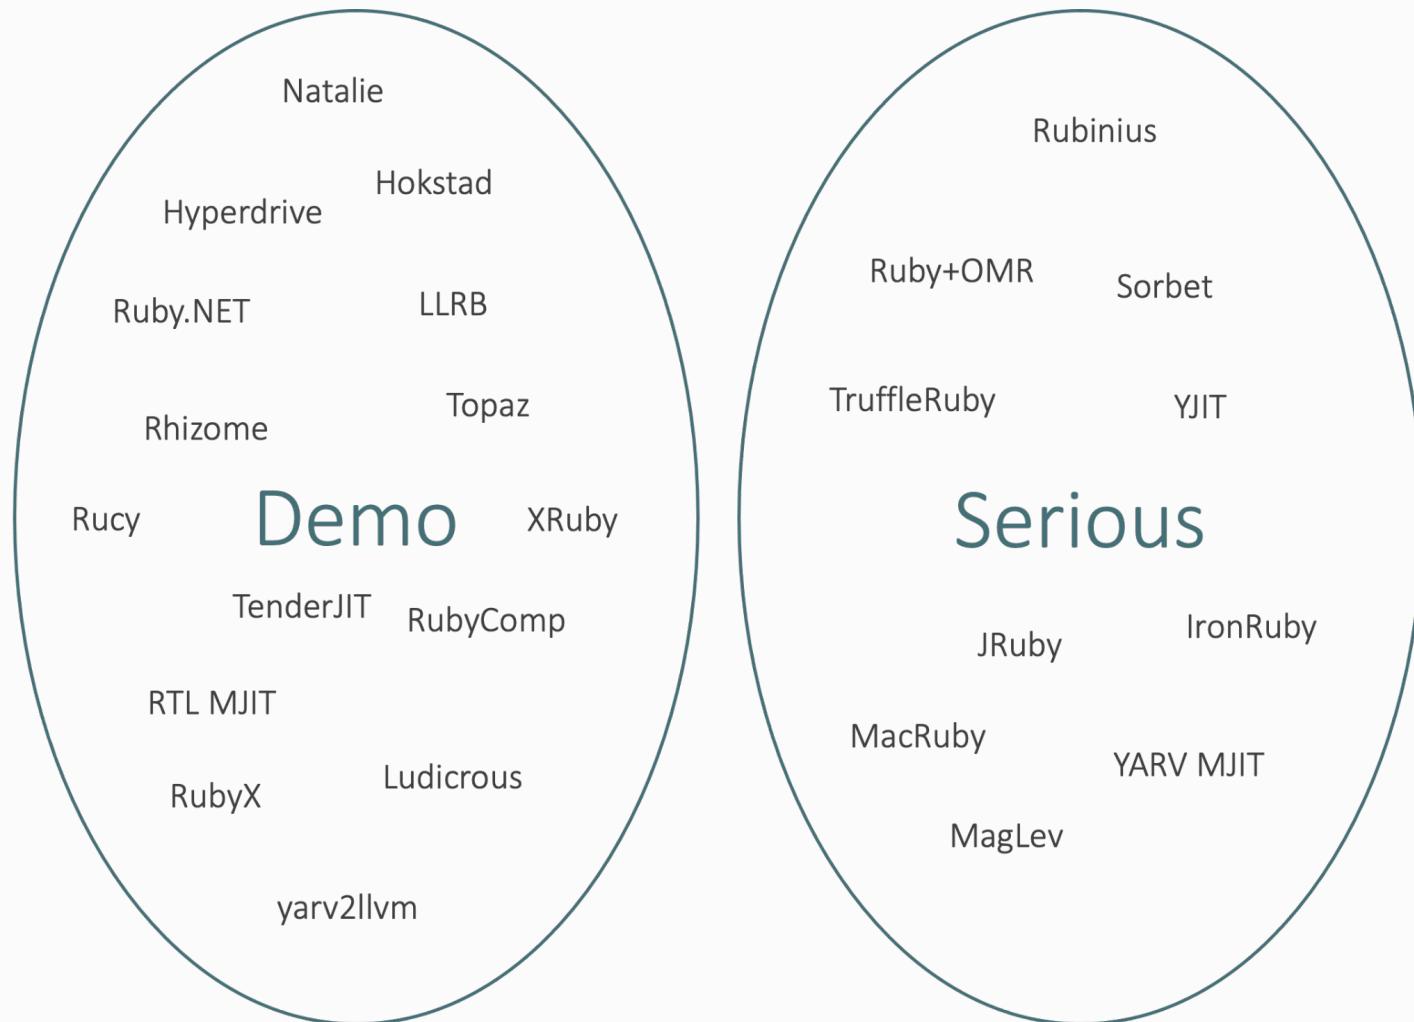


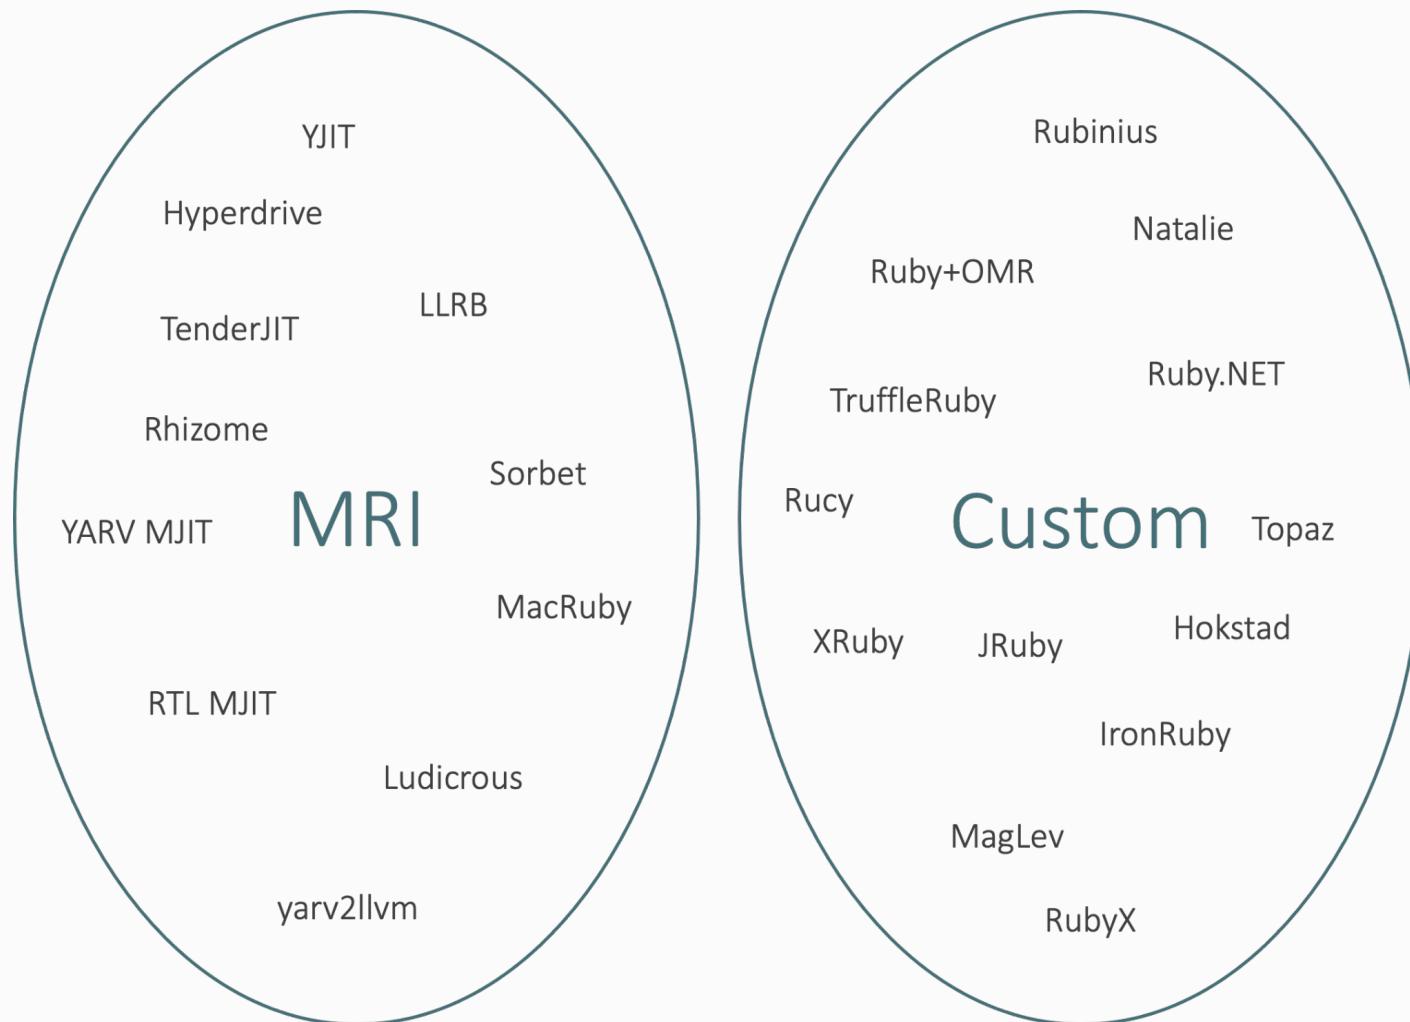


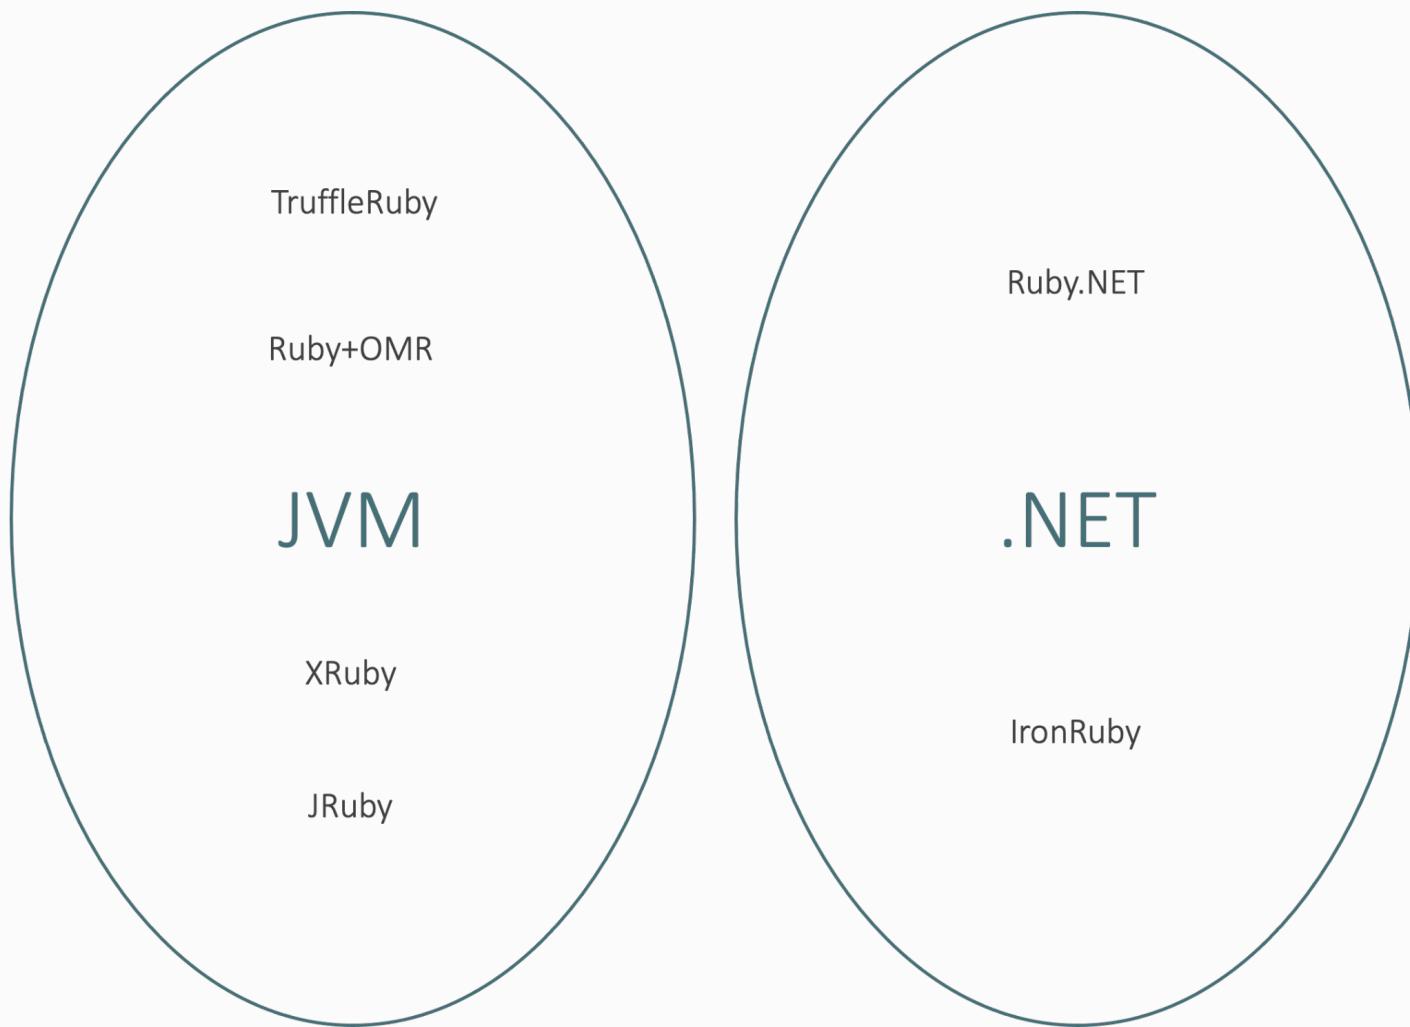






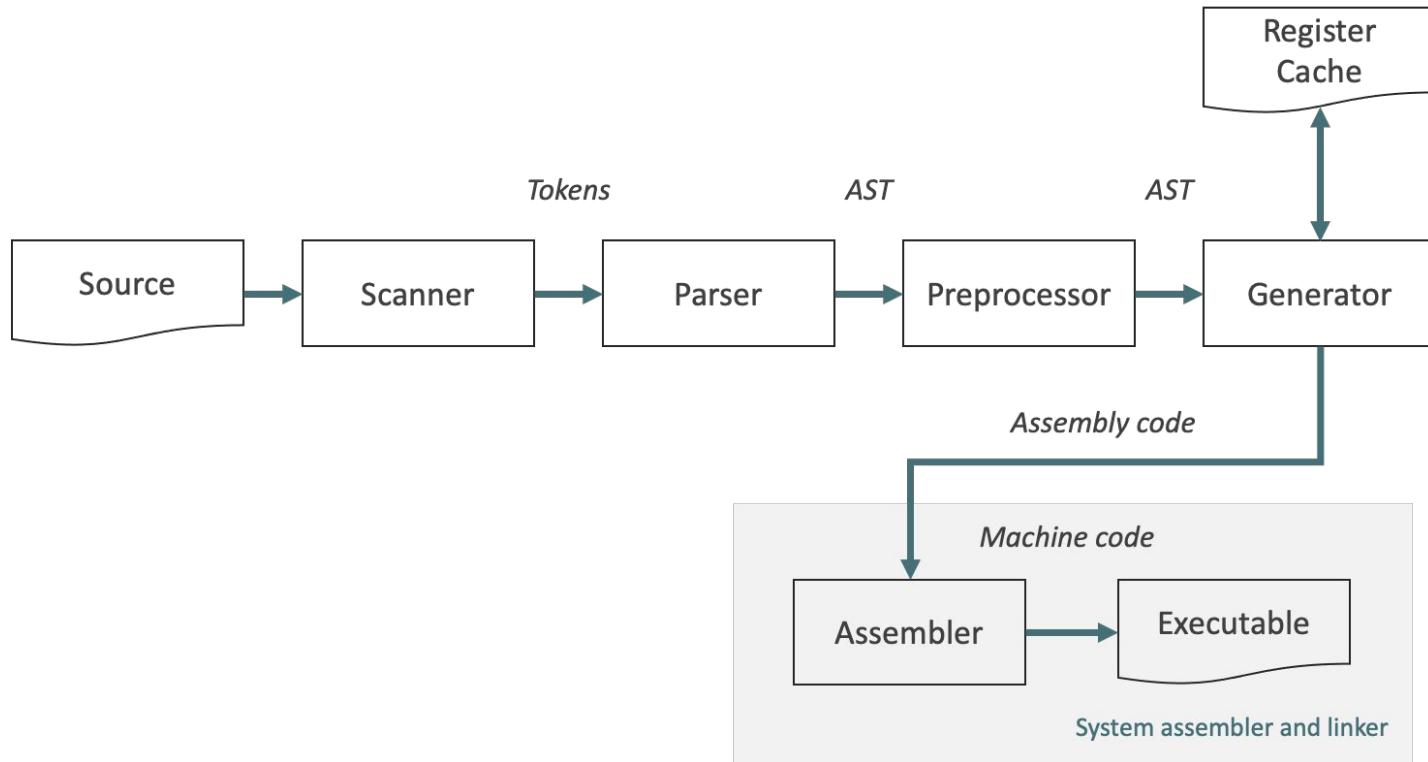






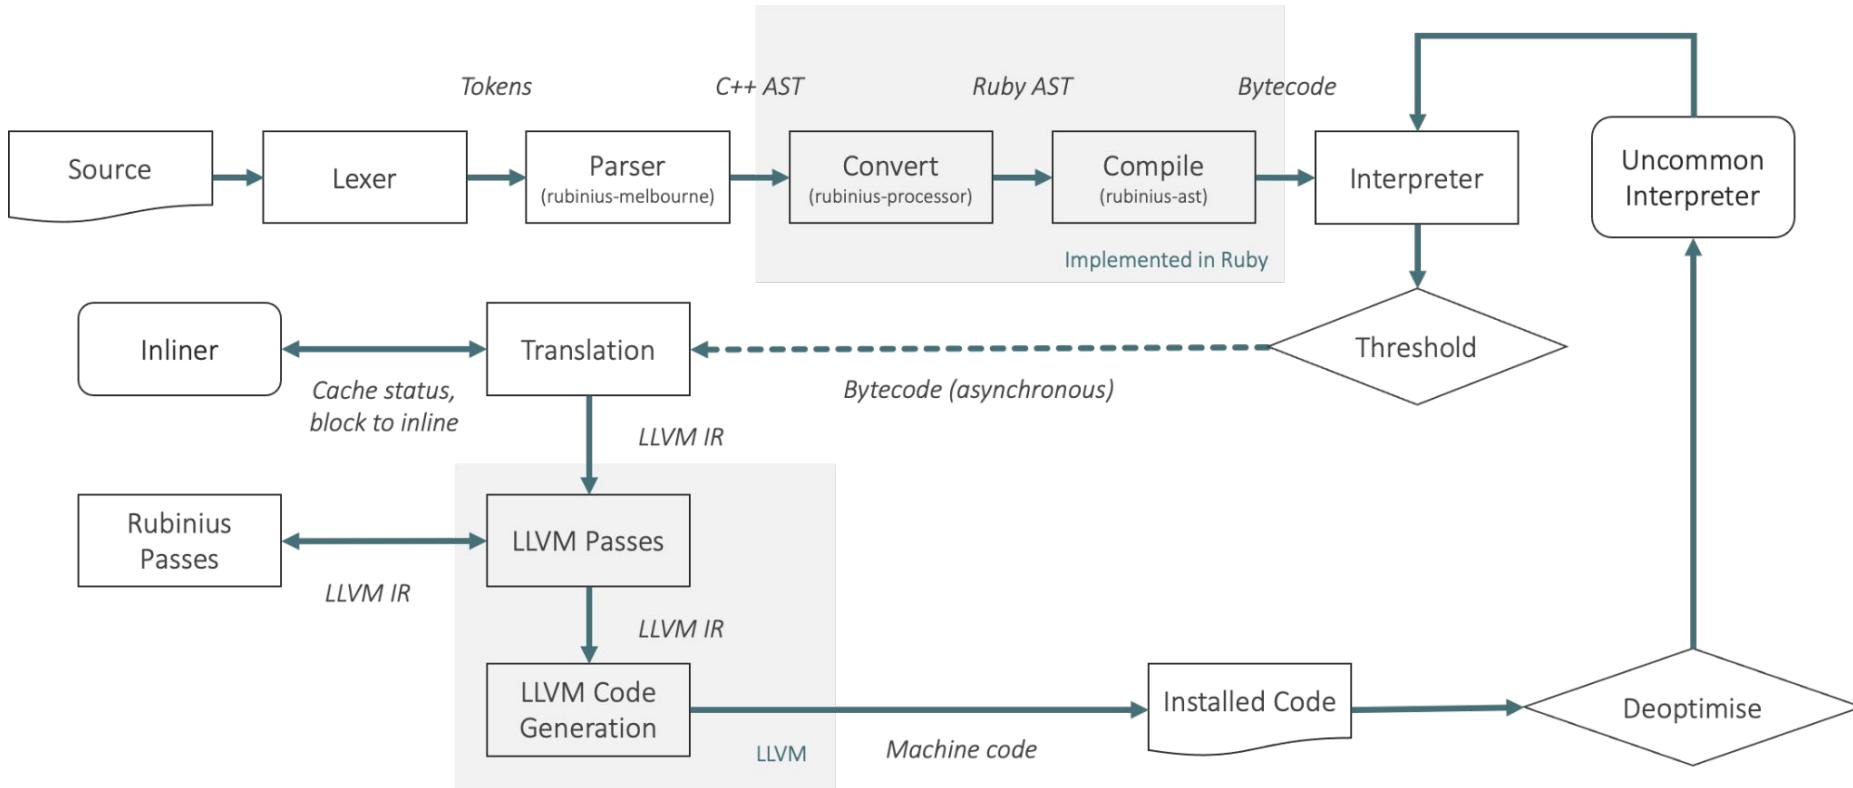
Deep dive into Hokstad





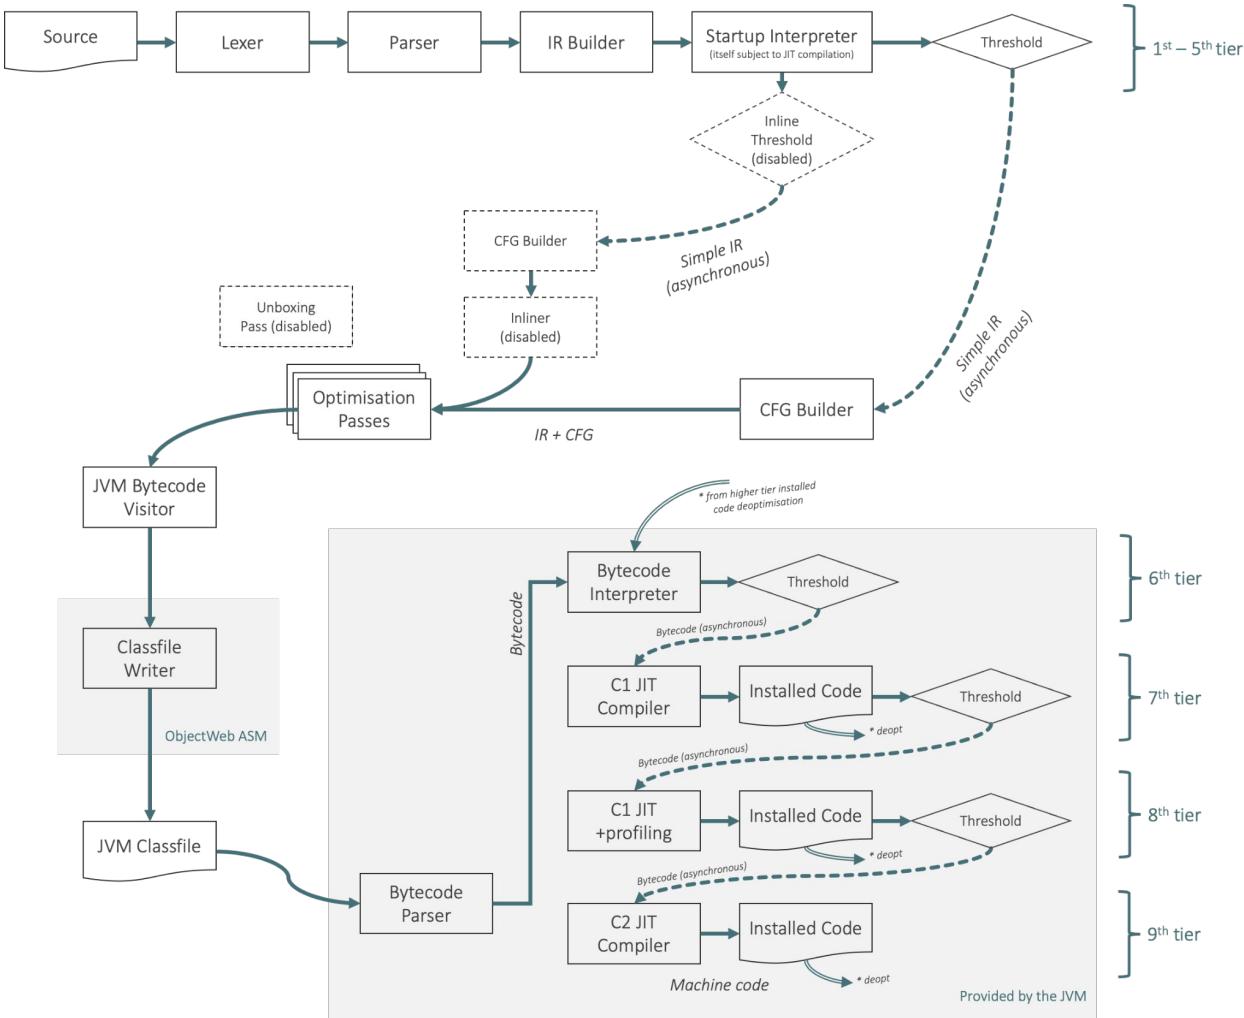
Deep dive into Rubinius

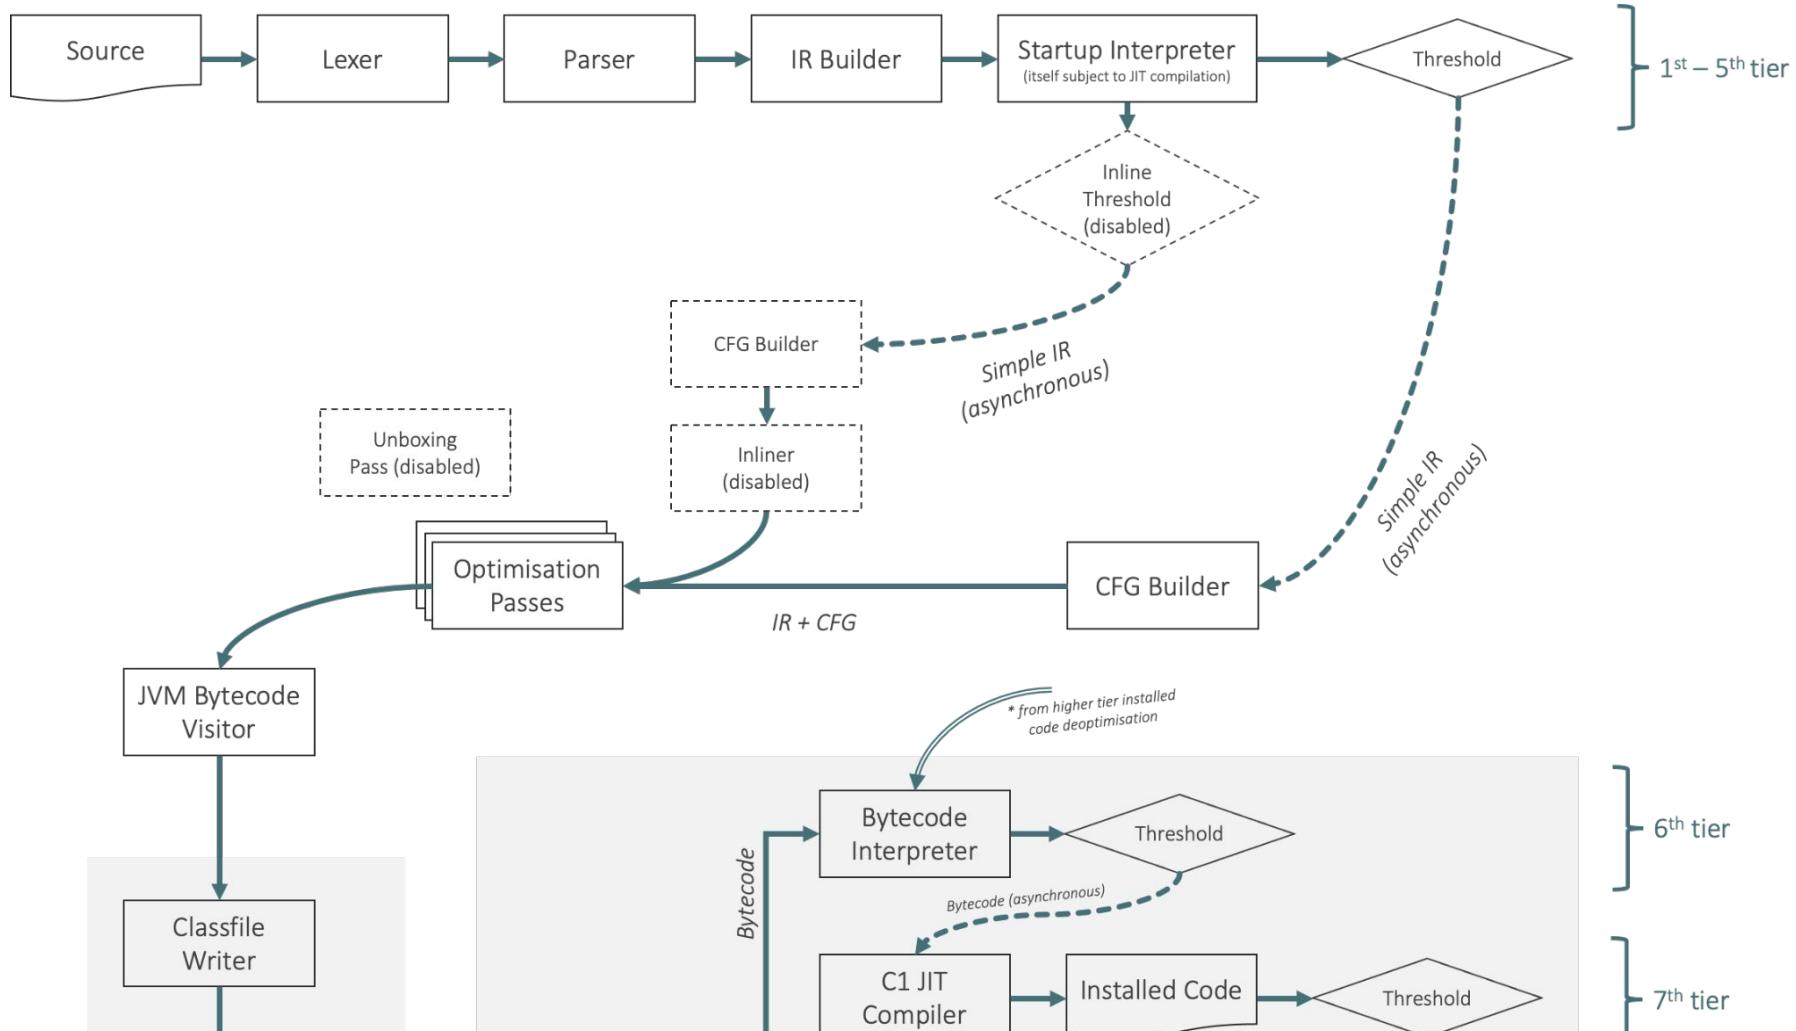


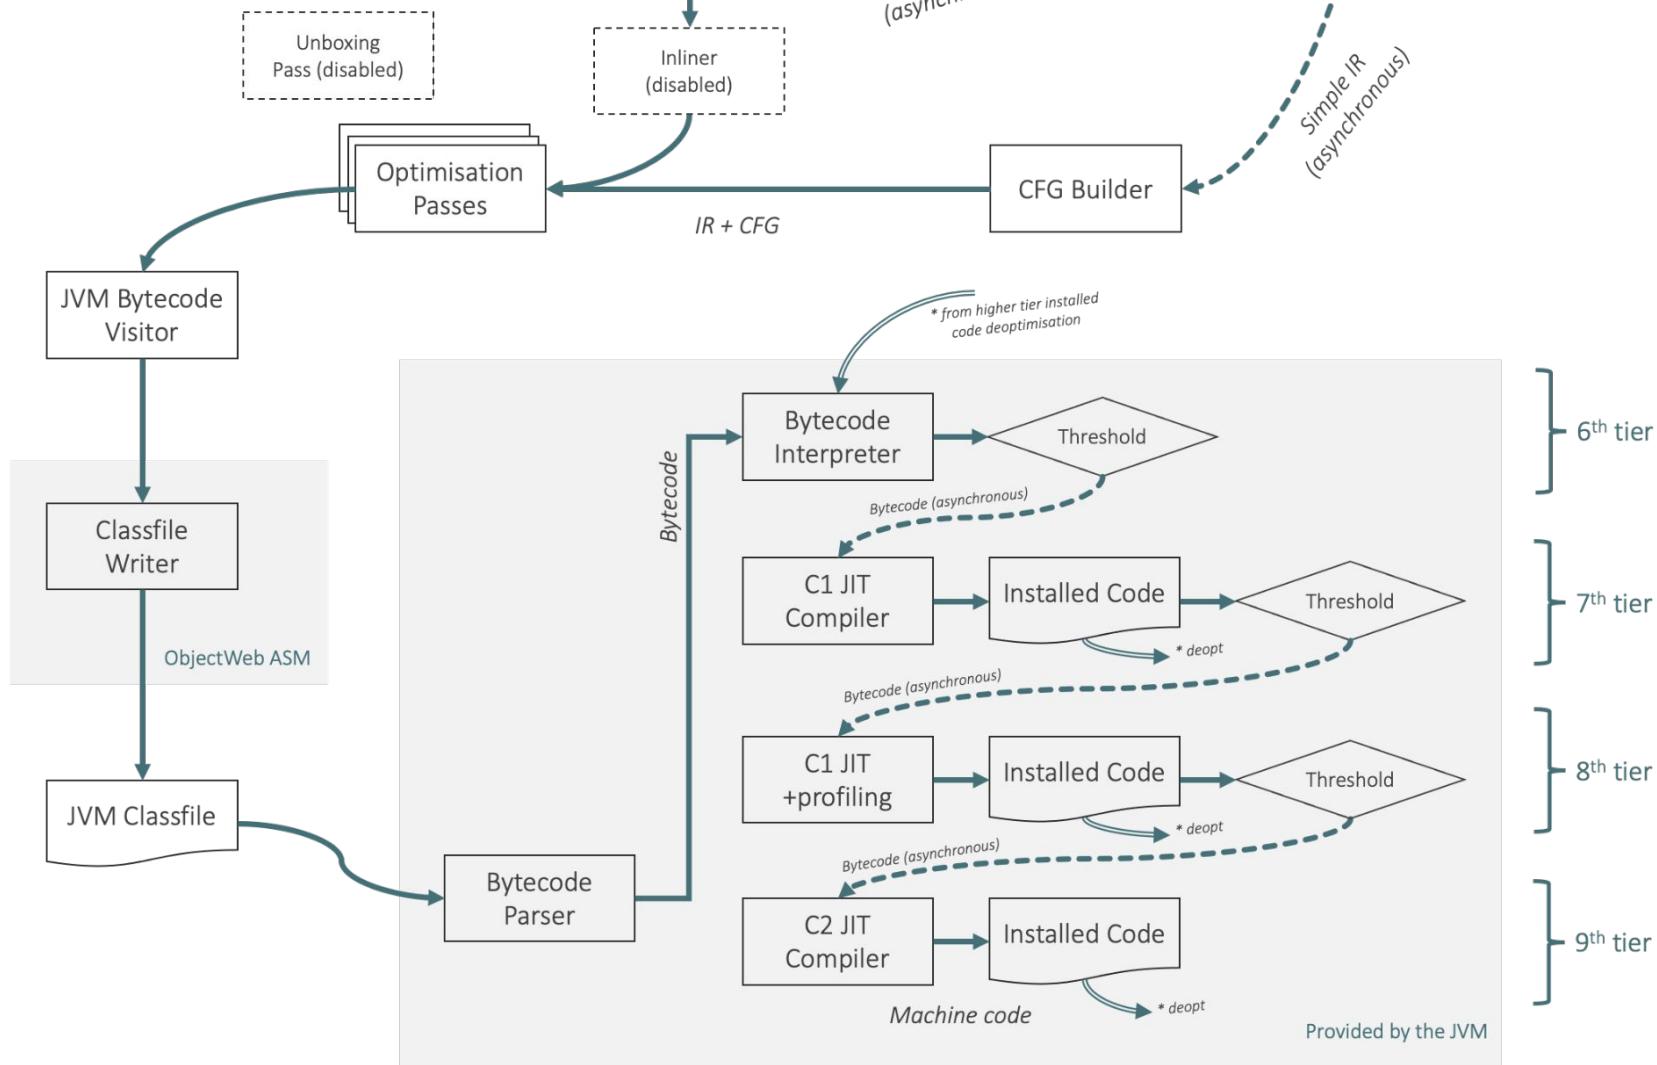


Deep dive into JRuby









Compiler research through Ruby

Tracing

Meta-tracing

Lazy Basic Block Versioning

Self-Specialising ASTs

Partial Evaluation and the Futamura Projection



Tracing

Research idea applied for Ruby in the Hyperdrive JIT by Jacob Matthews

Dynamic Native Optimization of Interpreters

Gregory T. Sullivan
gregs@ai.mit.edu
Artificial Intelligence
Lab

Derek L. Bruening
lyells@mit.edu
Laboratory for
Computer Science

Iris Baron
iris@csail.mit.edu
Laboratory for
Computer Science

Timothy Garnett
timothyg@csail.mit.edu
Laboratory for Computer Science

Saman Amarasinghe
sam@csail.mit.edu
Massachusetts Institute of Technology
Cambridge, MA

Abstract

For domain specific languages, “scripting languages”, dynamic languages, and for virtual machine-based languages, the most straightforward implementation strategy is to write an interpreter. A simple interpreter consists of a loop that fetches the next bytecode, dispatches it to a function, and then loops again. There are many ways to improve upon this simple mechanism, but as long as the execution of the program is driven by a representation of the program other than a stream of native instructions, there will be some “interpretive overhead”.

There is a long history of approaches to removing interpretive overhead from programming language implementations. In practice, what often happens is that, once an interpreted language becomes popular, pressure builds to improve performance until eventually a project is born that implements the language in native code as a compiler for the language. Implementing a JIT is usually a large effort, affects a significant part of the existing language implementation, and adds a significant amount of code and complexity to the overall code base.

In this paper, we present an innovative approach that dramatically removes much of the interpretive overhead from language implementations, with minimal instrumentation of the original interpreter. While it does not give performance improvements of a compiler, native compilers, our system provides an appealing point on the language implementation spectrum.

Categories and Subject Descriptors

D.3 [Software]: Programming Languages; D.3.4 [Programming Languages]: Processors—*Interpreters*

1 Introduction

Certain kinds of programming languages lend themselves to an interpreted implementation: languages with a high degree of dynamism (e.g., dynamic OO languages), or with a premium on fast startup and smaller programs (e.g., scripting languages), or where the performance of the interpreter is important. Indeed, if it can be argued that nearly all sufficiently complex applications have elements of interpretation throughout¹.

Typically, the language front end will transform the surface syntax of a program to an intermediate representation of the program (e.g., bytecode) that is then interpreted. As such, an interpreter consists of a loop that fetches the next bytecode, dispatches it to some routine handling that bytecode, then loops. There are many ways to improve upon this simple mechanism, but as long as the execution of the program is driven by a representation of the program other than a stream of native instructions, there will be some “interpretive overhead”.

We take an approach that is a combination of native *Just In Time* (JIT) compiler and *partial evaluation* techniques. We start with a dynamic optimization system called *DynamoRIO*, jointly developed at the labs and MIT. We give a quick overview of the DynamoRIO system in Section 2.

DynamoRIO records sequences of native instructions, which are subsequently partially evaluated with respect to the in-memory representation of the program being interpreted. Suppose the application program is represented as an immutable vector of bytecodes, and we have a long trace of native instructions. If we look carefully at the trace, there will typically be a value (we call it the “Logical PC”) used as an index into the bytecode vector, and there will be sequences of instructions as follows:

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee.
POM’10, June 12, 2010, San Diego, California, USA
Copyright 2010 ACM 1-4503-0090-0/10/06...\$10.00

¹ Greenspun’s Tenth Rule of Programming: “Any sufficiently complicated C or Fortran program contains an ad-hoc, informally-specified, bug-ridden, slow implementation of half of Common Lisp.”

Meta-tracing

Tracing the Meta-Level: PyPy's Tracing JIT Compiler

Carl Friedrich Bolz
University of Düsseldorf
STUPS Group
Germany
cfbolz@gmx.de

Maciej Fijalkowski
Inria & CNRS
fjal@merlinux.eu

Antonio Cuni
University of Genova
DISI
Italy
cuni@disi.unige.it

Armin Rigo
arigo@tunes.org

ABSTRACT

We attempt to apply the technique of Tracing JIT Compilers in the context of PyPy project, i.e., to programs that are interpreted by some dynamic languages, including Python. Tracing JIT compilers can greatly speed up programs that spend most of their time in loops in which they take similar code paths. However, applying an tracing JIT compiler to Python is hard, because the interpreter results in very limited or no speedup. In this paper we show how to guide tracing JIT compilers to greatly improve the speed of bytecode interpreters. One crucial point is to untilize the many different loops that are present and used by the implementer of the bytecode interpreter. We evaluate our technique by applying it to two PyPy interpreters: one is a small example, and the other one is the full Python interpreter.

1. INTRODUCTION

Dynamic languages have seen a steady rise in popularity in recent years. JavaScript is increasingly being used to implement full-scale applications which run within web browsers whereas other dynamic languages (such as Ruby, Perl, Python, PHP) are used for the server side of many web sites, as well as in areas dedicated to the web.

One of the disadvantages of dynamic languages is the performance penalties they impose. Typically they are slower than statically typed languages. Even though there has been a lot of research on the performance of various dynamic languages (in the SELF project, to name just one example [18]), those techniques are not as widely used as one would expect. Many of the performance issues that are considered by traditional bytecode interpreters without any advanced implementation techniques like just-in-time compilation. There are a number of reasons for this. Most of them boil down to the inherent complexities of using compilation. Interpreters are simple to implement, under-

stand, extend and port whereas writing a just-in-time compiler is a more error-prone task that is made even harder by the different features of a language.

A recent approach to getting better performance for dynamic languages is that of tracing JIT compilers [16, 7]. Writing a tracing JIT compiler is relatively simple. It can be implemented in a few days and it is needed for a tracing interpreter takes over some of the functionality of the compiler and the machine code generation part can be simplified.

The PyPy project is trying to find approaches to generally accelerate dynamic languages written as Python.

PyPy is a Python implementation of Python, but has now extended

its goals to be generally useful for implementing other dy-

namic languages as well. The general approach is to implement an interpreter for the language as a subset of Python.

This subset is chosen in such a way that programs in it

can be compiled into various target environments, such as C/C++ or Java. The PyPy project is described

in more detail in Section 2.

In this paper we discuss ongoing work in the PyPy project to implement the performance of interpreters written in the subset of Python that is chosen. The approach is that of a tracing JIT compiler. Contrary to the tracing JITs for dynamic languages that currently exist, PyPy's tracing JIT operates on one level below the user program, i.e., on the level of the interpreter, as opposed to the execution of the user program. The fact that the program the tracing JIT compiles is in our case always the interpreter brings its own set of problems to the tracing JIT. This is explained in the approach to interpreters in Section 3. By this approach we hope to arrive at a JIT compiler that can be applied to a variety of dy-

namic languages, given an appropriate interpreter for each of them. This is not necessarily automatic but needs

a number of hints from the interpreter author, to help the tracing JIT. The details of how the process integrates the tracing JIT with PyPy will be explained in Section 4. This work is the first step. It has already produced some promising results, which we will discuss in Section 6.

The contributions of this paper are:

- Applying a tracing JIT compiler to an interpreter.
- Finding techniques for improving the generated code.

2. THE PYPY PROJECT

The PyPy project¹ [21, 5] is an environment where flex-

Permission to make digital or hard copies of all or part of this work for personal use or internal distribution is granted without fee, provided that copies are not made or distributed for commercial advantage and that copies bear this notice and the full citation on the first page. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee.
Copyright 2008 ACM X-XXXXXX-XX-X/XX/XX ... \$5.00
<http://codesspeak.net/pypy>

Research idea applied for Ruby in the Topaz JIT by Alex Gaynor et al

Lazy Basic Block Versioning

Research idea applied for Ruby in YJIT by
Maxime Chevalier-Boisvert et al

YJIT: A Basic Block Versioning JIT Compiler for CRuby

Maxime Chevalier-Boisvert
Shopify
Canada
maxime.chevalierboisvert@shopify.com

Noah Gibbs
Shopify
United Kingdom
noah.gibbs@shopify.com

Jean Boussier
Shopify
France
jean.boussier@shopify.com

Si Xing (Alan) Wu
Shopify
Canada
paper@alanwu.email

Aaron Patterson
Shopify
United States
aaron.patterson@shopify.com

Kevin Newton
Shopify
United States
kevin.newton@shopify.com

John Hawthorn
GitHub
Canada
john@hawthorn.email

Abstract

Ruby is a dynamically typed programming language with a large breadth of features which has grown in popularity with the rise of the modern web, and remains at the core of the most popular Ruby implementations. While the default Ruby interpreter is fast, it does not always yield performance improvements on real-world software. Attempts to independently reimplement the Ruby language, such as [Ruby and TruffleRuby] have shown improvements in performance, but often lag behind CRuby when it comes to supporting new additions to the language, which limits their adoption.

We introduce YJIT, a new JIT compiler built inside CRuby based on a Lazy Basic Block Versioning (LBBV) architecture. We show that while our compiler does not match the peak performance of TruffleRuby, it offers near-100% compatibility with existing Ruby code, impressively fast warmup, and speedups from 15% to 19% on sizeable benchmarks based on real-world software.

CCS Concepts: • Software and its engineering → Just-in-time compilers.

Keywords: just-in-time, compiler, dynamically typed, optimization, bytecode, ruby

ACM Reference Format:

Maxime Chevalier-Boisvert, Noah Gibbs, Jean Boussier, Si Xing (Alan) Wu, Aaron Patterson, Kevin Newton, and John Hawthorn. 2021. YJIT: A Basic Block Versioning JIT Compiler for CRuby. In *Proceedings of the 13th ACM SIGPLAN International Workshop on Virtual Machines and Intermediate Languages (VMIL '21), October 19, 2021, Chicago, IL, USA*. ACM, New York, NY, USA, 8 pages. <https://doi.org/10.1145/3486606.3486781>

1 Introduction

Ruby is a programming language featuring dynamic typing and garbage collection, as well as both object-oriented and functional programming constructs. The language, which takes inspiration from Perl and Smalltalk, has grown in popularity over the years, and remains at the core of the implementation of many widely-used websites.

The feature-rich and highly dynamic nature of Ruby make it challenging to optimize. CRuby, the official Ruby implementation, includes a JIT compiler known as MJIT, but despite impressive numbers on synthetic benchmarks, MJIT has not thus far, yielded speedups in production workloads. Large-scale enterprise users, such as Shopify and GitHub typically leave it disabled, relying on the CRuby interpreter only.

There are ongoing efforts, such as [Ruby and TruffleRuby], to produce alternative Ruby implementations. These have shown great peak performance numbers, but for a variety of reasons, they have not yet been adopted for general deployment. They are still very limited. The current situation is not better than of alternative reimplementations of other existing languages. TruffleRuby, JRuby, PyPy and LuaJIT are all impressive implementation efforts with attractive performance numbers, but there is often a gap in terms of compatibility with and

Self-Specialising ASTs

AST Specialisation and Partial Evaluation for Easy High-Performance Metaprogramming

Chris Seaton
Oracle Labs
chris.seaton@oracle.com

Abstract

The Ruby programming language has extensive metaprogramming functionality, while most other languages, the use of whose features is idiomatic, and much of the Ruby ecosystem uses metaprogramming operations in the inner loops of libraries and applications.

The foundational techniques to make most of these metaprogramming operations fast have been known since the 1990s, but the lack of tool support for their application in practice is difficult enough that they are not widely applied in existing implementations of Ruby and other similar languages.

The TruffleRuby framework for writing self-specialising AST interpreters and the Graal dynamic compiler have been designed to make it easy to develop high-performance implementations of languages. We have found that the tools they provide also make it dramatically easier to implement efficient metaprogramming. In this paper we show how metaprogramming patterns from Ruby, show that with Truffle and Graal their implementation can be easy, concise, elegant and highly performant, and highlight the key tools that were needed from them.

Categories and Subject Descriptors D.3.4 [Programming Languages]: Processors—Run-time environments

Keywords Virtual Machines, Interpreters, Truffle, Graal, Ruby, Java

1. Introduction

Ruby is an impure, object-oriented, dynamically typed programming language with late-bound dispatch. This contrasts with languages such as Smalltalk and Python. Ruby has extensive metaprogramming functionality that allows parts of program execution to be controlled and observed using runtime data rather than data fixed in the source program. Examples include making method calls using a name that is a dynamic value, proxying method calls, dynamic code eval-

uation, and access to instance variables and method activation frames.

Many languages include some of this metaprogramming functionality, but a key difference in the Ruby programming language is cultural. While other programming languages tend to discourage use of metaprogramming, Ruby embraces it. Many examples can be found in memory management in Ruby, such as that of an application's inner loop, rather than just in offline operations such as set up and testing.

Ruby+Truffle (Seaton 2015) is a new implementation of Ruby using the Truffle (Wirthmueller et al. 2013) framework for self-specialising AST interpreters, and the Truffle interpreter (Traver et al. 2013a), a dynamic compiler that can be used to compile Truffle's ASTs using partial evaluation. In Ruby+Truffle we have developed techniques to make metaprogramming patterns from Ruby efficient (Seaton et al. 2015; Martí et al. 2015; Seaton et al. 2016). Dolgov et al. (2016) show the techniques needed for that are in many cases incremental developments on techniques already used for conventional program operations, but we will suggest that in practice they become difficult enough to implement that they are left unoptimised, and that this is a small set of key tools that Truffle and Graal provide that make it much easier to implement efficient metaprogramming.

We are implementing Ruby's metaprogramming functionality in this paper, but many of the methods described are available in other languages such as Python and JavaScript.

1.1 Contributions

Previous publications have described in depth some aspects of implementing metaprogramming functionality in Ruby and other languages implemented using Truffle and Graal. This paper provides a survey of how all the main metaprogramming functionality in Ruby has been implemented, including both concise descriptions of techniques already presented in other publications, and techniques not yet described in the literature.

This paper makes the following research contributions:

- Identification of typical Ruby metaprogramming functionality and the patterns in which they are used.

Research idea applied for Ruby in
TruffleRuby by Chris Seaton et al

Partial Evaluation

Practical Partial Evaluation for High-Performance Dynamic Language Runtimes

Thomas Würthinger^{*} Christian Wimmer^{*} Christian Humer^{*} Andreas Wöß^{*}
Lukas Stadler^{*} Chris Seaton^{*} Gilles Duboscq^{*} Doug Simon^{*} Matthias Grimmer[†]
^{*}Oracle Labs [†]Institute for System Software, Johannes Kepler University Linz, Austria
(thomas.wuerthinger, christian.wimmer, christian.humer, andreas.woess, lukas.stadler, chris.seaton,
gilles.m.duboscq, doug.simon)@oracle.com matthias.grimmer@jku.at

Abstract

Most high-performance dynamic language virtual machines duplicate language semantics in the interpreter, compiler, and runtime system. This violates the principle to not repeat yourself. In contrast, we define languages solely by writing an interpreter. The interpreter can specialize code, e.g., augments the interpreted program with type information and profiling information. Compiled code is derived automatically using partial evaluation techniques from the interpreter. This makes partial evaluation practical in the context of dynamic languages. It reduces the size of the compiled code and it allows all parts of an application to specialize when they are relevant for a particular part. If a specialization fails, execution transfers back to the interpreter, the program re-specializes in the interpreter, and later partial evaluation again transforms the new state of the interpreter to compiled code. We evaluate our approach against current implementations of JavaScript, Ruby, and R with best-in-class specialized production implementations. Our general-purpose compiler translates common benchmarks with production systems that have heavily optimized code for the one language they support. For our set of benchmarks, our speedup relative to the V8 JavaScript VM is 0.83x, relative to Ruby is 3.8x, and relative to GNU R is 5x.

CCS Concepts • Software and its engineering → Runtime environments

Keywords dynamic languages; virtual machine; language implementation; optimization; partial evaluation

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without the provider's explicit permission. For other uses, such as teaching materials, comments, or reprinting, please contact ACM's Copyrights and Permissions Department at www.acm.org. Copyright is held by the author(s). Publication rights licensed to ACM.
PERMISIÓN PARA HACER COPIAS DIGITALES O EN FORMA FÍSICA DE TODA O PARTE DE ESTA OBRA PARA USO PERSONAL O EN CLASES SE OTORGARÁ SIN EXPRESA AUTORIZACIÓN DEL PROPIETARIO. PARA OTROS USOS, COMO MATERIALES PARA PROFESORES, COMENTARIOS O REIMPRESIONES, POR FAVOR CONTACTE AL DÉPARTAMENTO DE DERECHOS Y PERMISOS DE ACM EN www.acm.org. EL DERECHO DE AUTOR SE MANTIENE EN LOS AUTORES. LOS DERECHOS DE PUBLICACIÓN SE OTORGAN A ACM.

DOI <https://doi.org/10.1145/3032311.3062801>

1. Introduction

High-performance virtual machines (VMs) such as the Java HotSpot VM or the V8 JavaScript VM follow the design that was first implemented for the Scheme language [23], a multi-tier optimization system using static optimization and deoptimization. The first execution tier, usually the interpreter or fast-compiling baseline compiler, enables fast start-up. The second execution tier, a dynamic compiler generating optimized code for interpreted code, provides good peak performance. Deoptimization transitions execution from the second tier back to the first tier, i.e., replaces stack frames of optimized code with frames of unoptimized code when an assumption made by the dynamic compiler no longer holds (see Section 5 for details).

Multiple tiers increase the implementation and maintenance costs for a VM. In addition to a language-specific optimizer, each tier needs to be implemented and tested separately. Even though the complexity of an interpreter or a baseline compiler is lower than the complexity of an optimizing compiler, implementing them is far from trivial. Additionally, the code needs to be maintained and ported to new architectures. But more importantly, the semantics of the language need to be implemented multiple times in different styles. For example, for a host language optimizer, it usually implements the compiler as interpreted code. For the second tier, language operations are implemented as graphs of language-specific compiler intermediate representations. And for the third tier, the code called from the interpreter (compiled code using remote calls), language operations are implemented in C or C++.

We propose to make the language interpreter written in a simple form as a language interpreter written in a managed high-level host language. Optimized compiled code is derived from the interpreter using partial evaluation. This approach and its main benefits were described in 1971 by G.J. Plotkin [13]. And for the first time, we show how to apply this approach to high-performance dynamic language projects. To the best of our knowledge no prior high-performance language implementation used this approach.

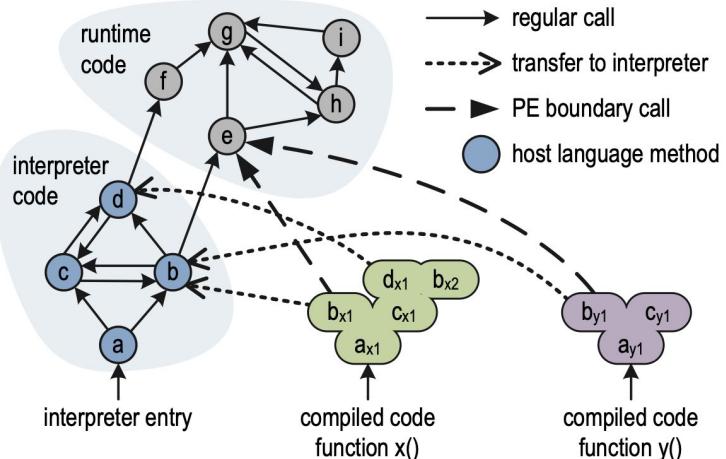


Figure 1: Interactions of interpreter code, compiled code, and runtime code.

Partial Evaluation



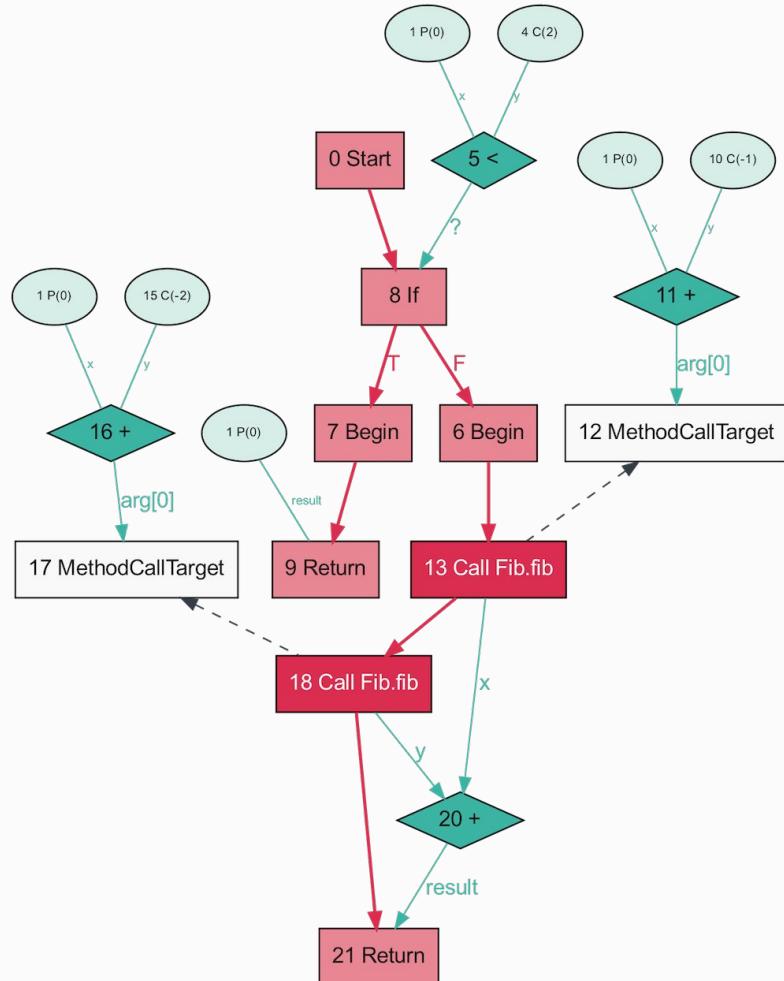
Some themes to ruminate on...



Where we might need to go next?

Intermediate representations

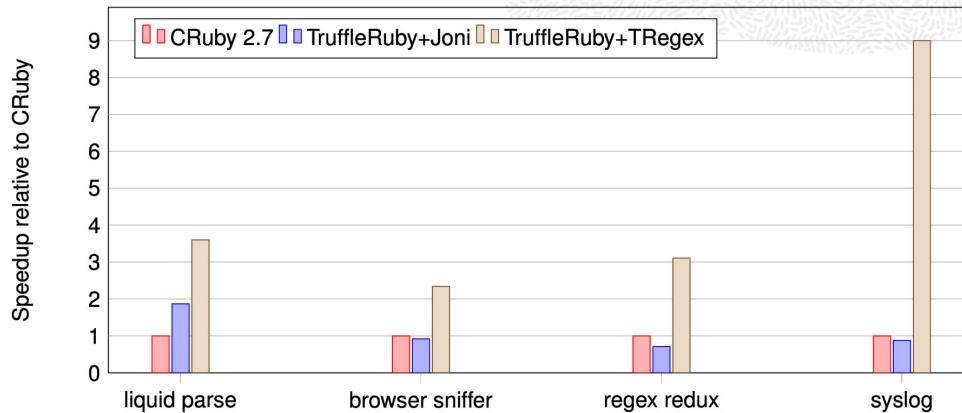




Where we might need to go next?

Polyglot compilation

Larger Regexp Benchmarks



Talks to catch on compilers at RubyConf

YJIT	Building a new JIT Compiler inside CRuby, Wed 10:50 am
JRuby	JRuby in 2021, Tue 1:40 pm
YARV MJIT	Optimizing Ruby Production Perf with MRI JIT, <i>watch anytime</i>
TruffleRuby	JIT Compiling Ruby Regexp, <i>watch anytime</i>
Sorbet Compiler	Compiling Ruby to Native Code with Sorbet and LLVM, Wed 1 pm
TenderJIT	Some Assembly Required, Tue 10:50 am



People to watch and read for Ruby compilers



Tim Morgan

Live-streaming his Natalie Ruby compiler on YouTube

youtube.com/c/TimMorgan



Vidar Hokstad 45-article blog post series about his Ruby compiler

hokstad.com/compiler



Aaron Patterson

Live-streaming his TenderJIT compiler on YouTube

youtube.com/c/TenderlovesCoolStuff

Why are Ruby compilers so cool?

- You can build one yourself
- So much freedom for how to design and architect them
- Tons of open space for new ideas
- You can apply research
- Really active area!



ruby-compilers.com

The Ruby Compiler Survey



rubybib.org

Academic writing on the Ruby programming language



@ChrisGSeaton

